

Machine Learning for Named Entity Recognition and Classification Report, Theoretical Component, and NERC-Research Preparation

Vera Langeberg

Abstract

Fill in a short abstract for your final submission.

We recommend aiming for a report of around 6. This does not include acknowledgements or references (they may appear on page 7 and later if needed). You can use appendices if you want to share more details. Below, we provide length indications per section. Please try to stick to them, but don't worry about receiving a lower grade if you don't quite manage. They are supposed to help you and prevent you from doing extra work.

1 Introduction

Length: $\pm 1/2$ - max 1 column for Introduction

Write a short introduction to your approach for the final submission. The introduction should include:

- A brief description of the task (not detailed)
- An outline of the ML approaches included in the report
- A brief summary of your results

Assignment 3

2 Task and Data

Length: ± 1.5 -2 columns for Task and Data.

2.1 Task

The exact task of NER given the CoNLL dataset is to identify and classify named entity spans in text into predefined categories. These categories are: locations (LOC), organizations (ORG), miscellaneous names (MISC), and persons (PER). Named entities are typically spans of text that can consist of one or more tokens. However, to simplify the task and transform it from span identification

into a token classification problem, the BIO (Beginning, Inside, Outside) tagging scheme is used. This scheme allows us to represent multi-token entity spans while still classifying individual tokens. Before them, there is a B- or I- to indicate whether it is the beginning or inside of an entity. The O label indicates that the token is outside of any named entity. This results in the following labels present in the dataset: "B-LOC", "I-ORG", "I-MISC", "B-MISC", "O", "B-ORG", "I-PER", "I-LOC", "B-PER"

For example, in the phrase "Vera Langeberg lives in Amsterdam":

The span "Vera Langeberg" is a person entity: "Vera" would be labeled as "B-PER" (beginning of person span) "Langeberg" would be labeled as "I-PER" (inside of person span) "lives" and "in" would be labeled as "O" (outside any entity span) The span "Amsterdam" is a location entity: "Amsterdam" would be labeled as "B-LOC" (beginning of location span) This approach allows the NER system to identify both single-token and multi-token named entity spans while performing token-level classification, effectively bridging the gap between span identification and token classification.

Assignment 2 Update if necessary

Assignment 3 Update if necessary

2.2 Dataset and Distribution

The dataset used in this project is the CoNLL-2003 dataset, which is specifically designed for Named Entity Recognition (NER) tasks (Sang and Meulder, 2003). The CoNLL format consists of four columns separated by spaces. Empty rows indicate sentence boundaries. The columns are as follows:

- Word or token
- Part-of-speech tag
- Chunk tag
- Named Entity tag

(Sang and Meulder, 2003)[p.2]

For example, in the line:

LONDON NNP B-NP B-LOC

"London" is the token, "NNP" is the part-of-speech tag (indicating a proper noun), "B-NP" is the chunk tag (indicating that it is the beginning of a noun phrase), and "B-LOC" is the named entity tag (indicating that it is the beginning of a location entity). The total number of tokens in the CoNLL training set is 203621. As can be seen in 1, the majority of the tokens are labeled with the O tag, indicating that they are outside of any named entity. The most frequent named entity label is B-LOC (7140), followed by B-PER (6600) and B-ORG (6321), as shown in 2. The least represented named entity label is I-MISC together with I-LOC. They only have 1155 and 1157 occurrences respectively in the training set.




Figure 1: Label distribution in the CoNLL training set.

In 9, the most common words for location, persons, organizations and miscellaneous names are shown. Here the IOB tags have been removed for clarity. For location, "u.s" is most common, "cup" is most common for miscellaneous names, "clinton" for persons and "of" for organizations.

Assignment 2 update if necessary

Assignment 3 update if necessary




Figure 2: Label distribution in the CoNLL training set (excluding "O").

2.3 Preprocessing

The only preprocessing that was done for the initial dataset analysis was lowercasing the tokens for the most frequent words per category. This was only done to collect the most frequent words, to avoid having the same word in different cases. Since, the case of the token is an important feature for NER, the original casing was kept for the model training. No stopwords removal or stemming/lemmatization was applied. Also no special characters were removed yet. This might be done in later assignments.

Assignment 2 update (if applicable)

Assignment 3 update (if applicable)

2.4 Evaluation Metrics

This paper will focus on token-level evaluation for the assignments. This means, that the BIO tags will be taken into account when calculating the metrics and have to match exactly. If only a partial match is made (so either wrong BIO tag, or wrong NE label), it will be considered incorrect.

Assignment 2 and 3: Update if necessary

3 Models and Features

Length: ± 2 columns for Models and Features.

3.1 Models

Assignment 1 The model used in the basic system is a Logistic Regression classifier from the sklearn library. Despite its name, it is mostly used for classification tasks. It has a linear model as its core that uses the sigmoid function to output probabilities for each

class. The model coefficients represent the importance of each feature for each class. The model was trained using the default parameters from sklearn on the training set.

Assignment 2 Alternative Methods (SVM, NB) The models that were used in addition to Logistic Regression are Support Vector Machines (SVM) and Naive Bayes (NB). SVM is a powerful classification algorithm that works by finding the optimal hyperplane that separates different classes in the feature space. It is particularly effective for high-dimensional data and can handle non-linear decision boundaries using kernel functions. In this assignment, the SVM model was implemented using the sklearn library with default parameters.

Naive Bayes is a probabilistic classifier based on Bayes' theorem, which assumes independence between features. It is computationally efficient and works well for text classification tasks, especially when the feature space is large. The NB model was also implemented using the sklearn library with default parameters.

Assignment 3 More advanced models (fine-tuning BERT, CRF for ReMA students only)

3.2 Features

Assignment 1 Feature exploration In this paper the features that were used in the basic system are:

- Token itself: The actual word provides crucial semantic information. Certain words are more likely to be part of specific entity types (e.g., "Inc." often indicates an organization).
- POS tags: Part-of-speech information helps identify potential named entities. For instance, proper nouns (NNP) are often good candidates for named entities.
- Capitalization: In English, this is a strong indicator for named entities, as proper nouns are typically capitalized. It helps distinguish between common nouns and potential named entities, though it can be less reliable at the beginning of sentences.

These features provide complementary information that aids in identifying and classifying named entities. The token itself carries semantic content, POS tags offer syntactic clues, and capitalization provides orthographic hints that are particularly useful for recognizing proper nouns.

The distribution of all these features can be found in the [A](#).

Assignment 2 More features

Assignment 3 More advanced features

4 Experiments and Results

Length: ± 3 columns (including figures and tables) for Experiments and Results. Additional material can be provided in the Appendix.

4.1 Evaluation

Assignment 1: first results The first results of the basic system on the development set are shown in figure 1 and the confusion matrix 3.

Label	Precision	Recall	F1-score	Support
B-LOC	0.870	0.746	0.803	1837
B-MISC	0.782	0.713	0.746	922
B-ORG	0.772	0.650	0.706	1341
B-PER	0.480	0.839	0.611	1842
I-LOC	0.758	0.525	0.621	257
I-MISC	0.836	0.457	0.591	346
I-ORG	0.640	0.465	0.539	751
I-PER	0.724	0.405	0.519	1307
O	0.989	0.993	0.991	42759
Accuracy			0.936	51362
Macro avg	0.761	0.644	0.681	51362
Weighted avg	0.943	0.936	0.936	51362

Table 1: Classification report with precision, recall, f1-score, and support for each label.

As can be seen, the model performs best on the B-LOC label with a 0.99 accuracy. The model struggles the most with the B-ORG label, with an accuracy of 0.40. The largest confusion happens between B-ORG that gets predicted as B-PER a lot. The macro precision is 0.761 and the macro recall is 0.644. The accuracy is with 0.936 quite high already.

Assignment 2: update

Assignment 3: update




Figure 3: Confusion matrix of the Logistic Regression model on the development set.

4.2 Feature Ablation

Assignment 3: feature ablation for one system

Note: if you decide after your feature ablation that you would like to discuss this before the Evaluation, you may move this subsection.

5 Error Analysis

Length: ± 1.5 -2 columns for Error Analysis.

Assignment 3: beyond the confusion matrix

6 Discussion

Length: ± 1 column for Discussion and Conclusion together.

Assignment 3

7 Conclusion

Assignment 3

Acknowledgements

References

- David Nadeau and Satoshi Sekine. 2007. [A survey of named entity recognition and classification](#). *Linguistic Investigations*, 30:3–26.
- Preslav Nakov, Zornitsa Kozareva, Alan Ritter, Sara Rosenthal, Veselin Stoyanov, and Theresa Wilson. 2019. [Semeval-2013 task 2: Sentiment analysis in twitter](#). *CoRR*, abs/1912.06806.
- Erik F. Tjong Kim Sang and Fien De Meulder. 2003. [Introduction to the conll-2003 shared task: Language-independent named entity recognition](#).

Warto, Supriadi Rustad, Guruh Fajar Shidik, Edi Noersasongko, Purwanto, Muljono, and De Rosal Ignatius Moses Setiadi. 2024. [Systematic literature review on named entity recognition: Approach, method, and application](#). *Statistics, Optimization & Information Computing*.

Theoretical Component

Assignment 1 - What is Machine learning?

My answer on the question "What is Machine learning" starts with explaining the intuition behind it. Machine learning is about machines learning, but this sounds still very silly and unclear. Therefore, let's first dive into the machines part. People should not think about machines like the first typewriter, or motor engines, or even your normal kitchen blender. We should also not think about the computers or calculators that run simple algorithms. Algorithms are mostly strictly defined, like the ones embedded in your regular calculators. These machines do not learn, they just follow a set of rules that are already defined. This is not machine learning, it is following a recipe. If the recipe is bad, you will get a burned cake. If you don't adjust the oven temperature next time, the cake will burn over and over again. The machines we are looking at are different. They can learn from data, identify patterns, and make decisions based on that data. To get back to our cake, this means that after one time burning the cake, it will lower the temperature next time, or even adjust the baking time. It can be undercooked that time, but it will learn from that as well, and adjust again. This learning from data and identifying patterns is the essence of machine learning.

When looking at it in a bit more formal way, machine learning is a subset of artificial intelligence (AI) that focuses on the development of algorithms and statistical models that enable computers to perform specific tasks without explicit instructions. Instead of being explicitly programmed to perform a task, machine learning algorithms use data to learn patterns and make predictions or decisions. This is done by training the algorithms on large datasets, allowing them to learn from the data and improve their performance over time. Machine learning can be categorized into three main types: supervised learning, unsupervised learning, and reinforcement learning. In supervised learning, the algorithm is trained on labeled data, where the input data is paired with the correct output. The algorithm learns to map the input to the output and can make predictions on new, unseen data. In unsupervised learning, the algorithm is trained on unlabeled data and must find patterns and relationships in the data on its own. Reinforcement learning involves training an agent to make decisions in an environment by rewarding or punishing it based on its actions. Key in the development of good machine learn-

ing algorithms are examples, features, and labels. Examples are the individual data points that the algorithm learns from. Features are the attributes or characteristics of the examples that are used to make predictions. Labels are the correct outputs or categories that the algorithm is trying to predict. To test the performance of machine learning algorithms, evaluation metrics such as accuracy, precision, recall, and F1-score are used. The evaluation is typically done on a separate dataset that was not used during training, known as the test set. A golden testset is a testset that is carefully curated and annotated by experts to ensure high quality and reliability.

Assignment 2

1.A.1: Task definition

A task definition of Subtask B (Nakov et al., 2019) with regards to input and desired output can be formulated as follows. Create a system that, given a message (single, unprocessed text message) from Twitter as input, decide the dominant sentiment from positive, negative, or neutral. For messages conveying both a positive and a negative sentiment, whichever is the stronger one was to be chosen. If no clear sentiment is present, neutral should be chosen. The output of the system are the predicted sentiment labels corresponding to the input messages.

1.A.2: Key Features for Sentiment Analysis in Tweets

To successfully determine the overall sentiment (positive, negative, or neutral) of a social media message, a system can rely on several types of numerical features that capture both the literal meaning of the words and the highly expressive style of online communication. The most fundamental features come from sentiment lexicons and N-grams. Lexicon-based features involve simply counting how many words in the tweet are listed in pre-compiled dictionaries as inherently positive (like "awesome") or negative (like "terrible"). This gives the system a strong initial signal of the message's overall leanings. N-grams, on the other hand, look at sequences of words (like "not nice" or "hate it") to capture context. This is crucial for recognizing phrases where one word negates the sentiment of another, ensuring the system doesn't misclassify "not nice" as a positive expression.

Beyond the words themselves, social media posts frequently use special elements to intensify

emotion. Features that track punctuation and emoticons can be highly valuable. Users often signal strong feelings using multiple exclamation points ("!!!") or specific emoticons like ":D" for joy or ":(" for sadness. Similarly, elongation and capitalization act as indicators of emphasis or anger. For example, a system can easily detect heightened emotion when a word is spelled with repeated letters ("sooo") or written entirely in capital letters ("SO ANGRY").

Finally, the system should incorporate features specific to the social media platforms. These can be hashtags or @-symbols that indicate towards which named entity the message is directed.

1.A.3: Feature representations

For sentiment analysis to get the desired outputs it starts by converting a text message into a Feature Vector ($\mathbf{x}$) within a Vector Space Model (VSM). The system first defines a large vocabulary of possible features, such as specific unigrams, the bigram "not nice," or the binary feature Has Positive Emoticon. Every message is then represented by an array of numbers where each component, x_j , records the magnitude or presence (e.g., count or a binary 1/0) of the corresponding feature j in that specific Tweet. For capitalization it would be useful to look at the ratio of capitalized words compared to not capitalized. This would also be a numerical representation. Important, the computer system sees only will see a high-dimensional, probably sparse (mostly zero-filled), numerical vector, not the human-readable text. The machine learning classifier, such as a Support Vector Machine (SVM), then utilizes these vectors to learn sentiment. During training, the classifier assigns a weight (w_j) to every feature j , with the weight indicating the feature's predictive power. For instance, the feature Count of Positive Words will receive a high positive weight, while the feature Bigram: "not nice" will receive a significant negative weight, reflecting its counter-intuitive sentiment. When classifying a new, unseen message, the system calculates a score by taking the linear combination of the feature vector and the learned weights, $\text{Score} = \sum_j w_j x_j$. This numerical score is then compared against established thresholds to assign the final class (positive, negative, or neutral). This mechanism reveals the core difference between human and machine perception. The human uses linguistic insight to understand the message's meaning, context, and intent (e.g., "This message is positive because the word 'brilliant' is

a strong compliment"). In contrast, the computer only performs mathematical algebra on the numbers, learning that the presence of the feature associated with "brilliant" has a specific high positive weight, which pushes the aggregate score above the positive threshold. The system's success is therefore entirely dependent on the quality of the initial human-engineered features that capture the necessary linguistic cues, so this part of the assignment is very important to me.

Assignment 3

Time spent

Week	Task	Time
1	Understand the assignment and getting everything setup	30 minutes
1	Watch the videos of module 1	2 hours
1	Finish 2.1 about understanding NER and CoNLL	1.5 hours
1	Work on 2.2.1 further exploration of the data	1 hour
1	Work on 2.2.2 feature hypothesis	1 hour
1	Started with 3.1, description of precision, recall etc	0.5 hour
2	Did all the week 2 tasks for assignment 1	2 hours
2	Incorporated the feedback for assignment 1	1.5 hours
3	Worked on the theoretical step of assignment 2	1 hour
3	Did a brief literature study for easy features	1.5 hours
3	Implemented the features from the literature study	2.5 hours
Total	15h	

Table 2: Time overview.

A Example Appendix

This is an appendix. You can include additional analyses and tables here (e.g. samples you analyzed for your error analysis).

NERC-research preparation

2.1.1 - NER given the CoNLL dataset

In this course we will use the CoNLL-2003 dataset, which is specifically designed for Named Entity Recognition (NER) tasks ([Sang and Meulder, 2003](#)). The exact task of NER given the CoNLL dataset is to identify and classify named entities in text into predefined categories. These categories are: locations (LOC), organizations (ORG), miscellaneous names (MISC), and persons (PER). Before them, there is a B- or I- to indicate whether it is the beginning or inside of an entity. The O label indicates that the token is outside of any named entity. This results in the following labels present in the dataset: "'B-LOC', 'I-ORG', 'I-MISC', 'B-MISC', 'O', 'B-ORG', 'I-PER', 'I-LOC', 'B-PER'".

The conventions that are followed in this dataset is IOB tagging scheme. This means that when a named entity consists of multiple tokens, the first token is labeled with a B- (beginning) and the subsequent tokens with I- (inside). For example, in the entity "New York City", "New" would be labeled as B-LOC, "York" as I-LOC, and "City" as I-LOC. There is consistency in the labeling, meaning that if a token is labeled as I-ORG, it must be preceded by either B-ORG or another I-ORG.

Punctuation marks are represented as separate tokens and are labeled with the O tag, indicating that they are outside of any named entity.

2.2.1 - How does CoNLL work?

The CoNLL format consists of four columns separated by spaces. Empty rows indicate sentence boundaries. The columns are as follows:

- Word or token
- Part-of-speech tag
- Chunk tag
- Named Entity tag

([Sang and Meulder, 2003](#))[p.2]

For example, in the line:

```
LONDON NNP B-NP B-LOC
```

"London" is the token, "NNP" is the part-of-speech tag (indicating a proper noun), "B-NP" is the chunk tag (indicating that it is the beginning of a noun phrase), and "B-LOC" is the named entity tag (indicating that it is the beginning of a location entity).

2.2.1 - Distribution of the labels

To investigate the distribution of the labels in the CoNLL dataset, I wrote a small piece of code that extracts the labels from the dataset and counts their frequency. The total number of tokens in the training set is 203621. The results show that the majority of the tokens are labeled with the O tag, indicating that they are outside of any named entity. This was what I expected, as not all tokens in a text are part of named entities. When looking at the named entity labels, I found that the most frequent named entity label is B-LOC (7140), followed by B-PER (6600) and B-ORG (6321). The least represented named entity label is I-MISC together with I-LOC. They only have 1155 and 1157 occurrences respectively in the training set.

2.2.2 - Hypothesis about linguistic features that might help NER

I hypothesize that the following linguistic and orthographic features might help improve NER performance:

- Part-of-speech tags: Certain parts of speech, such as proper nouns (NNP), are more likely to be named entities. Including POS tags as features can help the model identify potential named entities.
- Capitalization: Named entities often start with capital letters. Including a feature that indicates whether a token is capitalized can help the model distinguish named entities from common nouns.
- Contextual




Figure 4: Distribution of POS tags in the CoNLL training set.

words: The words surrounding a token can provide important context for identifying named entities. "the" or "an" or "a" is very common before a named entity. Including features that capture these patterns can help the model recognize named entities.

To test if these features have influence on the NER performance, I wrote some code in the notebook. Starting with the POS tags. I extracted the POS tags from the dataset and counted their frequency. The results show that the most common POS tags in the CoNLL training set are NNP 4. Next, I looked at the capitalization feature. I counted the number of capitalized tokens versus non-capitalized tokens. The results show that a significant (about 25%) portion of the tokens in the CoNLL training set are capitalized 5. Finally, I analyzed the contextual words preceding named entities. I extracted the tokens that appear before named entities and counted their frequency. The results show that certain tokens, such as "the", "of", and "in", are more common before named entities 6 and 7. This didn't correspond exactly with my hypothesis, as I expected more articles like "a" and "an" to be present.

Next, I explored the most common words per category. Since the figure 8 is a bit hard to read, I created a new graph without the IOB tags. This graph shows the most common words for location, persons, organizations and miscellaneous names. 9. For location, "u.s" is most common, "cup" is most common for miscellaneous names, "clinton" for persons and "of" for organizations.

3. 1. Confusion matrix

In this section I discuss the terms precision, recall and F-score in my own words. The precision of a model is the proportion of true predictions compared to all instances that got the true label. So how much that got filtered out, was actually correct. The recall of a model is the proportion of true predictions compared to all instances that actually have that label. How much of what should have been filtered out, was actually filtered out. The F-score is the harmonic mean of precision and recall, providing a single metric to evaluate the model's performance. This takes both precision and recall into account, and balances them. A scenario where high precision is most important is your spam email filter. We all hate when we missed an important email because it got filtered as spam. When you would have a high precision, everything that is in the spam folder is actually spam. A scenario where high recall is most important is in medical dermatology diagnosis. It is worse to miss a cancer diagnosis (false negative) than to have a false alarm (false positive). These are the classic examples that were also used in previous courses I followed.

3.2. Code for confusion matrix and metrics

I already implemented the code to calculate recall, precision and F1-score. There is also a way to plot the confusion matrix in there.




Figure 5: Capitalization distribution in the CoNLL training set.




Figure 6: Top 15 most common tokens before named entities in the CoNLL training set.




Figure 7: Top 15 most common tokens preceding named entities in the CoNLL training set.

3.3. Discuss token level evaluation and other related questions

In token-level evaluation, each token in the text is treated as an individual instance for evaluation purposes. In the case of the CoNLL dataset, each token is assigned a label indicating whether it is part of a named entity or not. And each datapoint in our dataset is a token. The advantages of analyzing at token level are:

- **Granularity:** Token-level evaluation provides a more fine-grained analysis of model performance, allowing us to identify specific tokens that are misclassified.
- **Simplicity:** Token-level evaluation is straightforward to implement and interpret, as it directly compares predicted labels to true labels for each token.

However, there are also disadvantages to token-level evaluation:

- **Overestimation of performance:** correctly classifying individual tokens may not necessarily translate to correctly identifying entire named entities, leading to an overestimation of model performance.
- **Not like the real-world:** In the real world, complete named entities are more important than individual tokens.

The advantages of using span-level evaluation instead of token-level evaluation are:

- **More realistic:** Span-level evaluation better reflects the real-world application of NER (it looks at the entire NE instead of individual tokens).
- **Better assessment of entity recognition:** Span-level evaluation provides a more accurate assessment of a model's ability to recognize complete named entities.

Since there are BIO tags present in the CoNLL dataset, it is possible to convert token-level predictions to span-level predictions. This can be done by grouping consecutive tokens with the same named entity label into spans. The spans will then be evaluated based on whether the entire span is correctly identified and classified as a named entity. No, BIO tagging is not strictly necessary for span-level evaluation, but it does facilitate the conversion from token-level to span-level predictions.

To deal with partial matches in span-level evaluation, we would need to use the strict evaluation approach. In this approach, a predicted span is considered correct only if it exactly matches the true span




Figure 8: Most common words per named entity label in the CoNLL training set.




Figure 9: Most common words per named entity label (without IOB) in the CoNLL training set.

in both boundaries and label. This means that if a predicted span overlaps with the true span but does not exactly match it, it is considered incorrect. This strict evaluation approach ensures that the model is evaluated based on its ability to accurately identify and classify complete named entities.

Since the option of evaluating at span-level is advanced, I will focus on token-level evaluation for the assignments. This means, that the BIO tags will be taken into account when calculating the metrics and have to match exactly. If only a partial match is made (so either wrong BIO tag, or wrong NE label), it will be considered incorrect. In the notebook, I already implemented this evaluation function.

Assignment 2

New features from literature study After conducting a literature study on easy-to-implement features for NER, I found several features that could potentially improve the performance of our NER model. These features include:

- previous predicted label
- presence of digits
- position in the sentence

(Warto et al., 2024) provides a comprehensive overview of various features used in NER tasks. From the figure on page 920 in the paper, I concluded that the most used features were word embeddings, POS-tagging and character embeddings. Since word embeddings are part of the next step, I focused on the POS-tagging and the preceding POS tag. In (Nadeau and Sekine, 2007), the authors discuss various features for NER, including word features and lexical features. The table on page 8 in the paper shows several word-level features like the presence of digit and the position of the token in the sentence.

I implemented the features above in the code and will evaluate their impact on the model's performance in the next section.